


# Data X


## ML Summary and Next Steps in Illustrations Data X: A Course on Data, Signals, and Systems

Ikhtlaq Sidhu  
Chief Scientist & Founding Director,  
Sutardja Center for Entrepreneurship & Technology  
IEOR Emerging Area Professor Award, UC Berkeley


## Overview



Supervised Learning  
Algorithms




Unsupervised Learning  
Algorithms



Semi-supervised  
Learning Algorithms

Data<sup>x</sup>



```
#Setting up for Supervised learning  
# First clean: use mapping + buckets
```


```
# X = matrix of data – e.g 1000 rows  
# Y = In sample responses
```

```
# Typically we want to split in to  
training data and test data
```

```
X_train = X[0:500]  
Y_train = Y[0:500]  
X_test = X[501:1000]  
Y_test = Y[501:1000]
```

Data<sup>X</sup>

## Linear Regression Illustration



```
#Setting Linear Regression in sklearn  
from sklearn import linear_model
```

```
model= linear_model.LinearRegression()  
model.fit(X_train, Y_train)
```

```
Y_pred_train = model.predict(X_train)
```

```
Y_pred_test = model.predict(X_test)
```

```
# Compare Y_pred_test with Y_test for  
error.
```

Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>

Data<sup>X</sup>

## Logistic Regression Illustration




Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>

Data<sup>x</sup>

## Logistic Regression Illustration

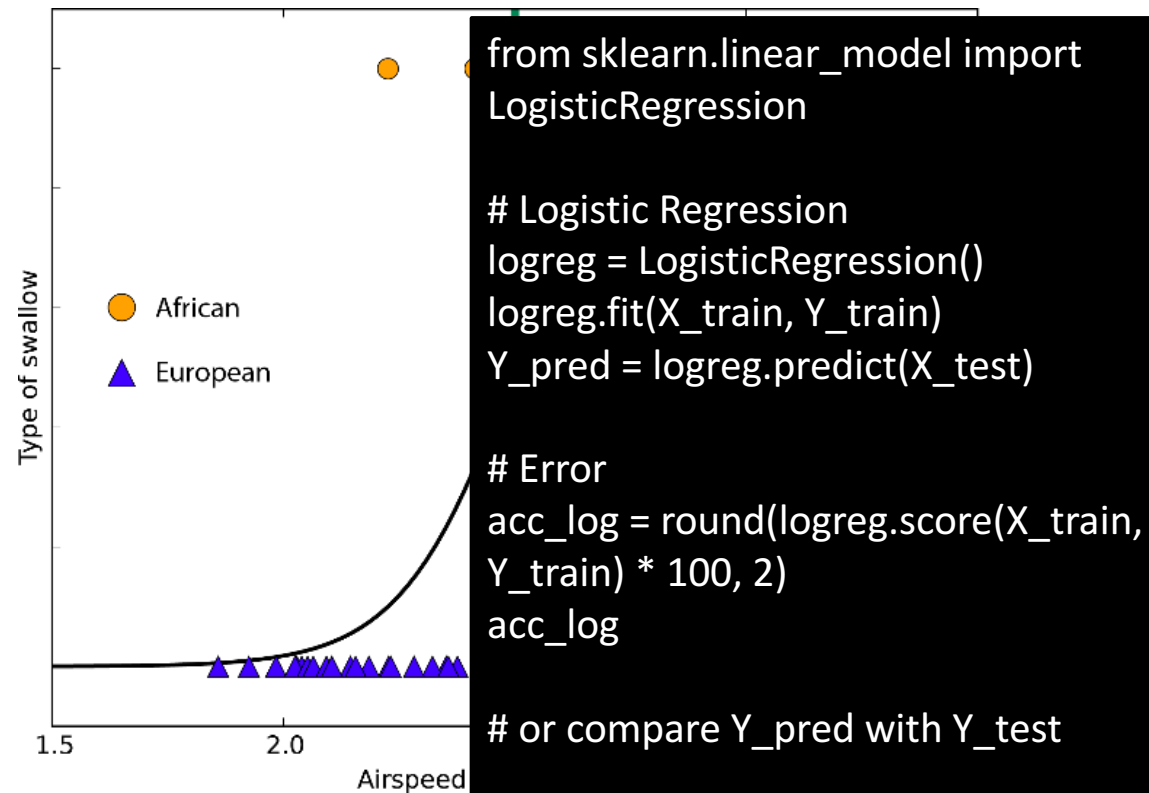



Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>

Data<sup>X</sup>



# Support Vector Machine (SVM) Illustration




*A typical support vector machine class boundary maximizes the margin separating two classes*

Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>



# Support Vector Machine (SVM) Illustration



*A typical support vector machine class boundary maximizes the margin separating two classes*

```
from sklearn.svm import SVC,  
LinearSVC
```

```
svc = SVC()  
svc.fit(X_train, Y_train)  
Y_pred = svc.predict(X_test)
```

```
# Error  
acc_svc = round(svc.score(X_train,  
Y_train) * 100, 2)  
acc_svc
```


```
# or compare Y_pred with Y_test
```

Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>

Data<sup>X</sup>



# Support Vector Machine (SVM) Illustration



*A typical support vector machine class boundary maximizes the margin between the two classes*

```
from sklearn.svm import SVC, LinearSVC
```

```
# Linear SVC
```

```
linear_svc = LinearSVC()  
linear_svc.fit(X_train, Y_train)
```

```
Y_pred = linear_svc.predict(X_test)
```

```
# Error:
```

```
acc_linear_svc =  
round(linear_svc.score(X_train, Y_train) *  
100, 2)acc_linear_svc
```


```
# or compare Y_pred with Y_test
```

Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>


Data<sup>X</sup>

# KNN / K Means Illustration


the data



NN classifier



5-NN classifier



**KNN Method:** Find the k nearest images and have them vote on the label (i.e. take the mode)





Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>


Data<sup>x</sup>

## K Means / KNN Illustration


the data



NN classifier



5-NN classifier



**KNN Method:** Find the k nearest images and have them vote on the label (i.e. take the mode)

Illustration Source: <https://www.kdnuggets.com/2015/05/knn.html>

```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier(n_neighbors = 3)
```

```
knn.fit(X_train, Y_train)
```

```
Y_pred = knn.predict(X_test)
```

```
acc_knn = round(knn.score(X_train, Y_train) * 100, 2)
```

```
acc_knn
```

```
# or compare Y_pred with Y_test
```

Data<sup>x</sup>

# Decision Tree Illustration






Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>

DataX

## Decision Tree Illustration



```
from sklearn import tree
```

```
decision_tree = DecisionTreeClassifier()  
decision_tree.fit(X_train, Y_train)  
Y_pred = decision_tree.predict(X_test)
```

```
# Error
```

```
acc_decision_tree =  
round(decision_tree.score(X_train, Y_train) *  
100, 2)  
acc_decision_tree
```

```
# or compare Y_pred with Y_test
```

Illustration Source: <https://docs.microsoft.com/en-us/machine-learning/sklearn/decision-tree/tutorial>

Data<sup>x</sup>

## Our experiment with the Titanic Data Set

|  | <b>Model</b>               | <b>Score</b> |
|--|----------------------------|--------------|
|  | Random Forest              | 86.76        |
|  | Decision Tree              | 86.76        |
|  | KNN                        | 84.74        |
|  | Support Vector Machines    | 83.84        |
|  | Logistic Regression        | 80.36        |
|  | Linear SVC                 | 79.01        |
|  | Perceptron                 | 78.00        |
|  | Naive Bayes                | 72.28        |
|  | Stochastic Gradient Decent | 72.28        |




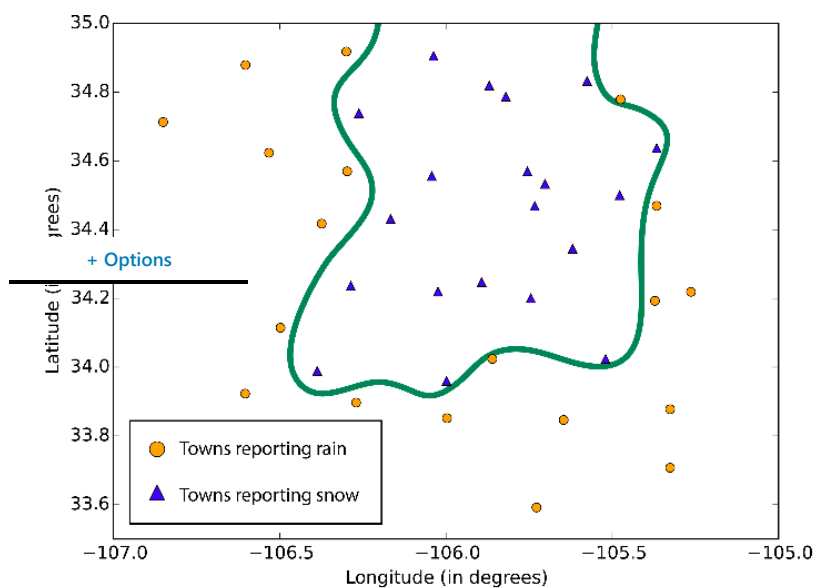
More Accuracy  
Generally more training time  
More risk of overfitting

Less Accuracy  
Generally less computation





# Neural Network Illustration

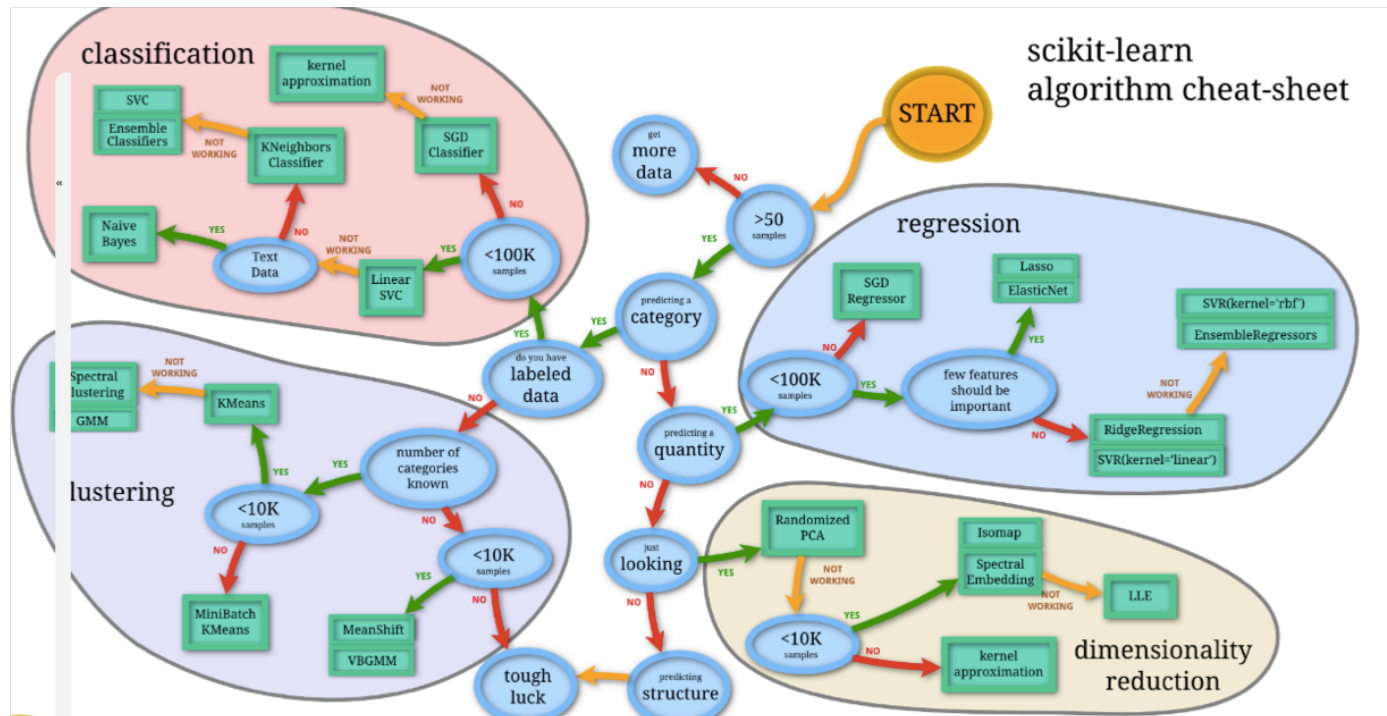


*The boundaries learned by neural networks can be complex and irregular*

Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>

DataX

# Scikit-Learn Algorithm



Data<sup>x</sup>

| Algorithm                           | Accuracy | Training time | Linearity | Parameters | Notes                                  |
|-------------------------------------|----------|---------------|-----------|------------|--|
| <b>Two-class classification</b>     |          |               |           |            |  |
| logistic regression                 |          | •             | •         | 5          |  |
| decision forest                     | •        | ○             |           | 6          |  |
| decision jungle                     | •        | ○             |           | 6          | Low memory footprint                   |
|                                     |          |               |           |            |  |
| boosted decision tree               | •        | ○             |           | 6          | Large memory footprint                 |
| neural network                      | •        |               |           | 9          | Additional customization is possible   |
|                                     |          |               |           |            |  |
| averaged perceptron                 | ○        | ○             | •         | 4          |  |
| support vector machine              |          | ○             | •         | 5          | Good for large feature sets            |
| locally deep support vector machine | ○        |               |           | 8          | Good for large feature sets            |
| Bayes' point machine                |          | ○             | •         | 3          |  |
| <b>Anomaly detection</b>            |          |               |           |            |  |
| support vector machine              | ○        | ○             |           | 2          | Especially good for large feature sets |
| PCA-based anomaly detection         |          | ○             | •         | 3          |  |
| K-means                             |          | ○             | •         | 4          | A clustering algorithm                 |

| Algorithm                         | Accuracy | Training time | Linearity | Parameters | Notes   |
|-----------------------------------|----------|---------------|-----------|------------|---|
| <b>Multi-class classification</b> |          |               |           |            |   |
| logistic regression               |          | •             | •         | 5          |   |
| decision forest                   | •        | ○             |           | 6          |   |
| decision jungle                   | •        | ○             |           | 6          | Low memory footprint                            |
| neural network                    | •        |               |           | 9          | Additional customization is possible            |
| one-v-all                         | -        | -             | -         | -          | See properties of the two-class method selected |
| <b>Regression</b>                 |          |               |           |            |   |
| linear                            |          | •             | •         | 4          |   |
| Bayesian linear                   |          | ○             | •         | 2          |   |
| decision forest                   | •        | ○             |           | 6          |   |
| boosted decision tree             | •        | ○             |           | 5          | Large memory footprint                          |
| fast forest quantile              | •        | ○             |           | 9          | Distributions rather than point predictions     |
| neural network                    | •        |               |           | 9          | Additional customization is possible            |
| Poisson                           |          |               | •         | 5          | Technically log-linear. For predicting counts   |
| ordinal                           |          |               |           | 0          | For predicting rank-ordering                    |

DataX

Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>

End of Section

Data<sup>x</sup>